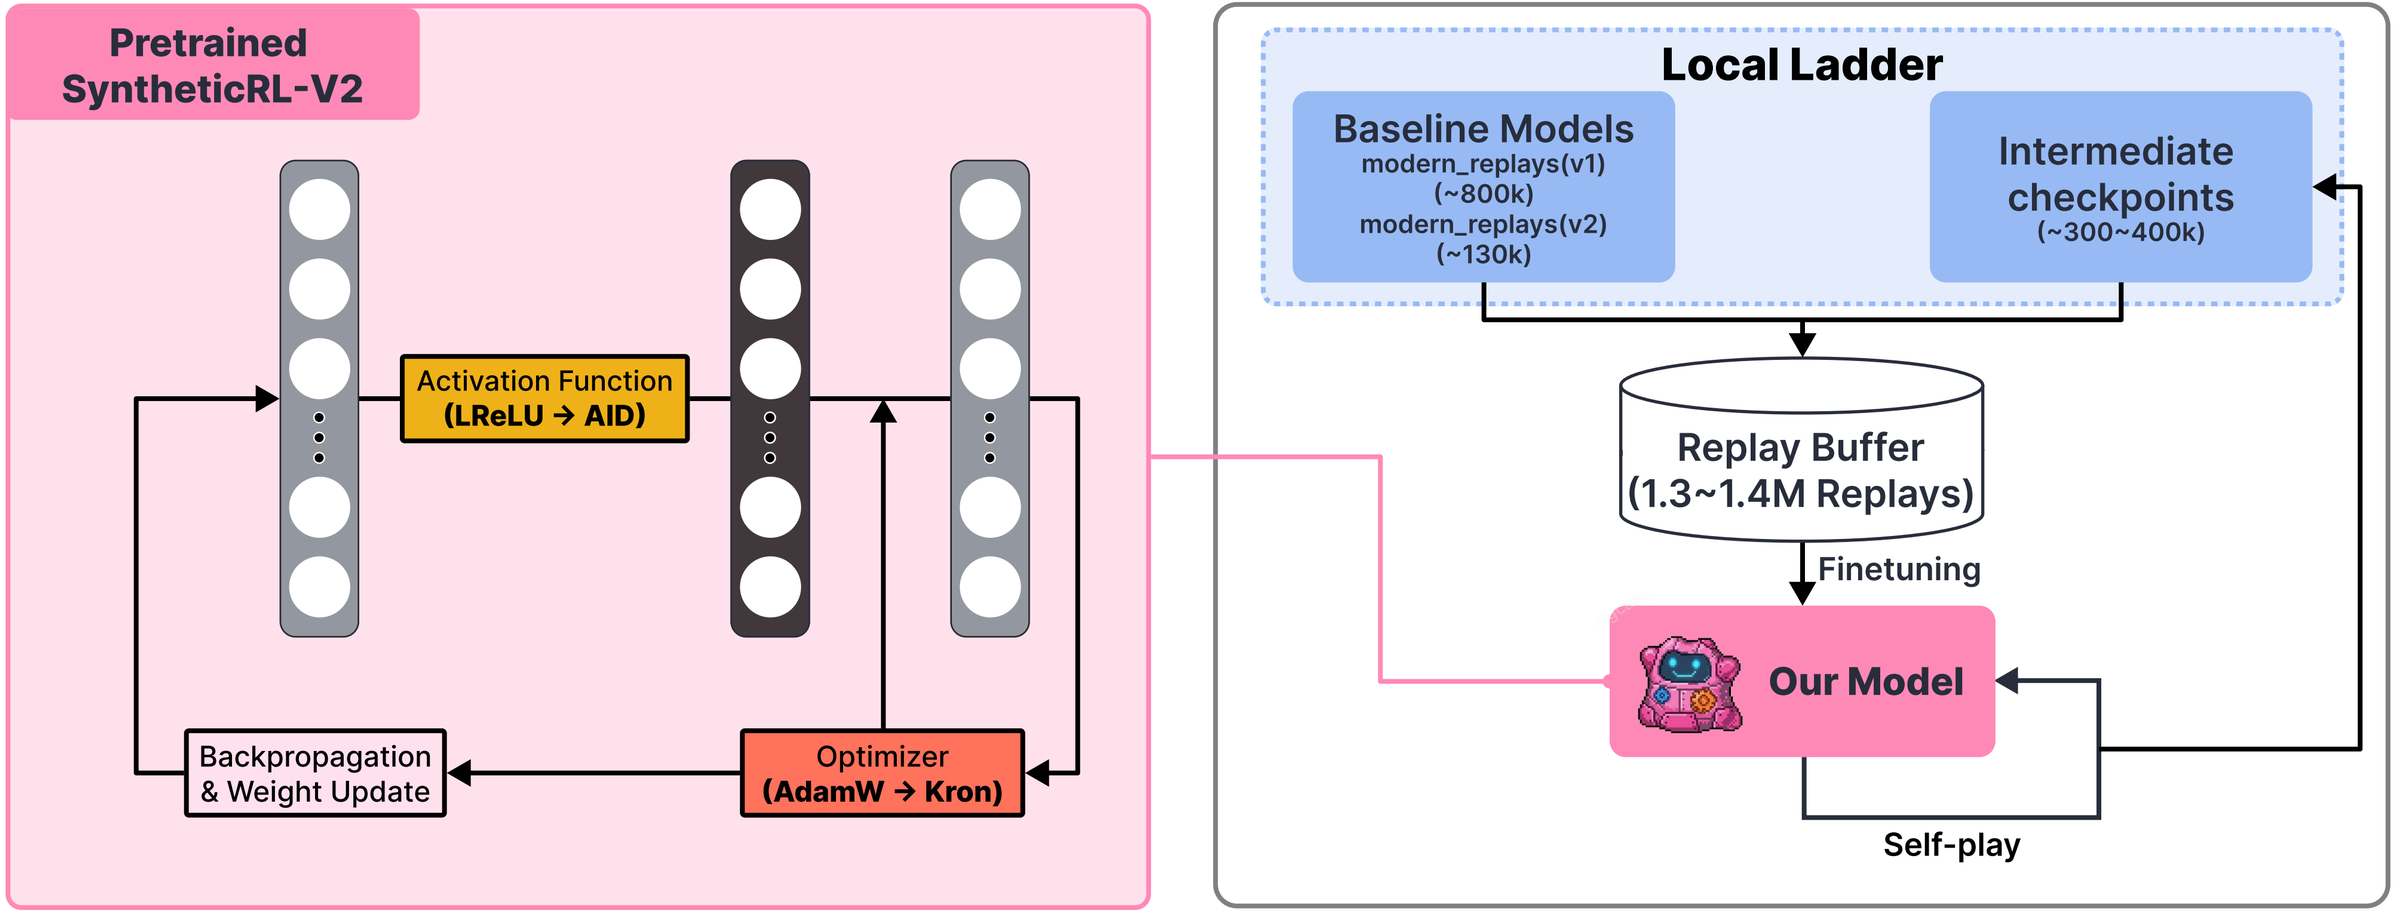
Pretrained SyntheticRL-V2
Activation Function
(LReLU → AID)
Backpropagation & Weight Update
Optimizer
(AdamW → Kron)
Local Ladder
Baseline Models
modern_replays(v1)
(~800k)
modern_replays(v2)
(~130k)
Intermediate checkpoints
(~300~400k)
Replay Buffer
(1.3~1.4M Replays)
Finetuning
Our Model
Self-play